

CS162  
Operating Systems and  
Systems Programming  
Lecture 7

Synchronization I:  
Concurrency, Mutual Exclusion, and Atomic Operations

September 19<sup>th</sup>, 2024

Prof. Ion Stoica

<http://cs162.eecs.Berkeley.edu>

# Recall: the Dispatch Loop

---

- Conceptually, the scheduling loop of the operating system looks as follows:

```
Loop {  
    RunThread();      // Needs to return to loop every now and then!  
    ChooseNextThread();  
    SaveStateOfCPU(curTCB);  
    LoadStateOfCPU(newTCB);  
}
```

- This is an *infinite* loop
  - One could argue that this is all that the OS does
- Should we ever exit this loop???
  - When would that be?

# Running a thread

---

Consider first portion: `RunThread()`

- How do I run a thread?
  - Load its state (registers, PC, stack pointer) into CPU
  - Load environment (virtual memory space, etc)
  - Jump to the PC
- Note: We give control of processor/core to user code!!
  - OS is *not running* because user code is running
- How does the OS get control back?
  - Internal events: thread returns control *voluntarily*
  - External events: thread gets *preempted*

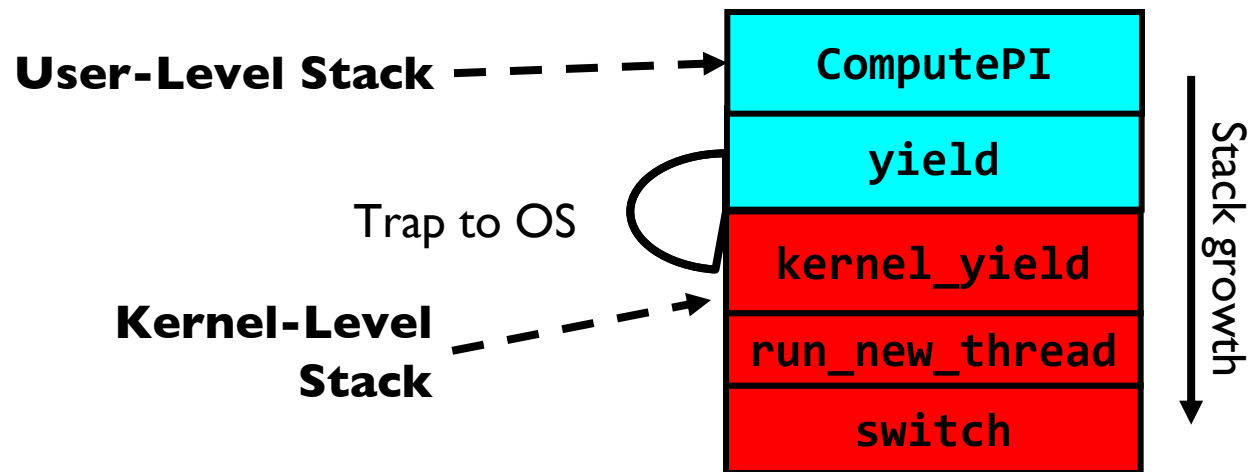
# Internal Events

---

- Blocking on I/O
  - The act of requesting I/O implicitly yields the CPU
- Waiting on a “signal” from other thread
  - Thread asks to wait and thus yields the CPU
- Thread executes a `yield()`
  - Thread volunteers to give up CPU

```
computePI () {  
    while (TRUE) {  
        ComputeNextDigit ();  
        yield ();  
    }  
}
```

# Stack for Yielding Thread



- How do we run a new thread?

```
run_new_thread() {  
    newThread = PickNewThread();  
    switch(curThread, newThread);  
    ThreadHouseKeeping(); /* Do any cleanup */  
}
```

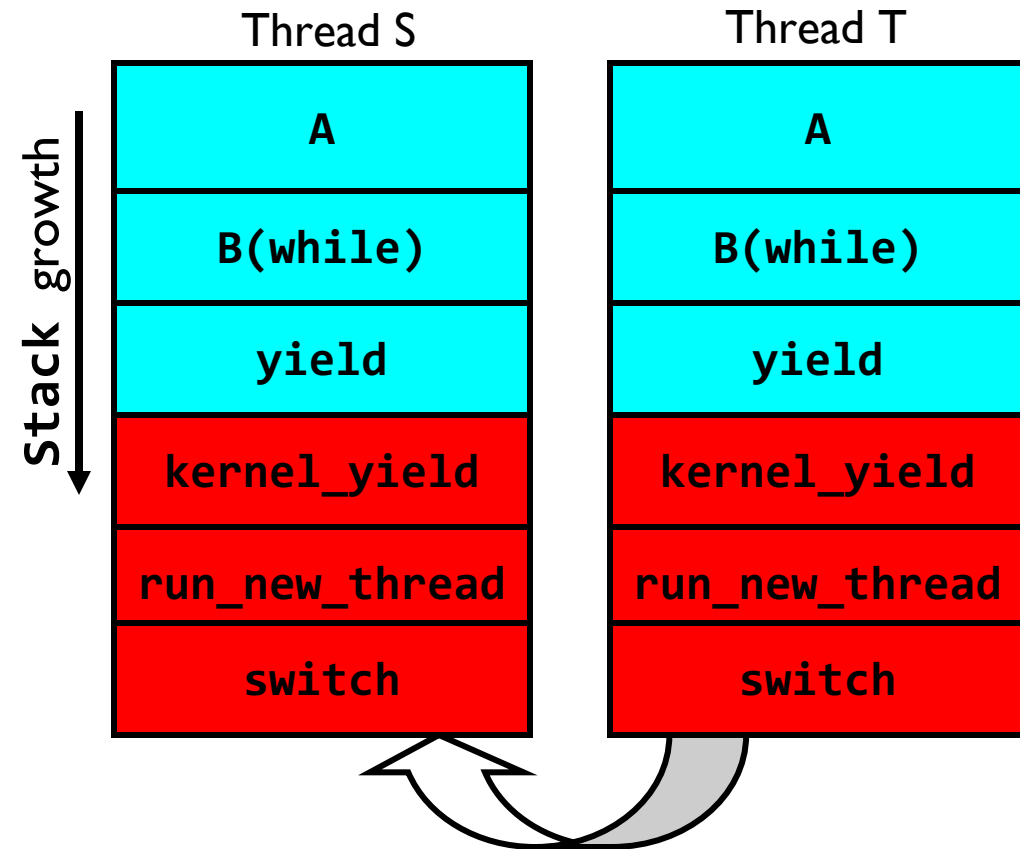
- How does dispatcher switch to a new thread?
  - Save anything next thread may trash: PC, regs, stack pointer
  - Maintain isolation for each thread

# Stacks for Yield with Multiple Threads

- Consider the following code blocks:

```
proc A() {  
    B();  
}  
proc B() {  
    while(TRUE) {  
        yield();  
    }  
}
```

- Suppose we have 2 threads:
  - Threads S and T
  - Assume that both have been running for a while



# Saving/Restoring state (often called “Context Switch”)

---

```
Switch(tCur, tNew) {  
    /* Unload old thread */  
    TCB[tCur].regs.r7 = CPU.r7;  
    ...  
    TCB[tCur].regs.r0 = CPU.r0;  
    TCB[tCur].regs.sp = CPU.sp;  
    TCB[tCur].regs.retpc = CPU.retpc; /*return addr*/  
  
    /* Load and execute new thread */  
    CPU.r7 = TCB[tNew].regs.r7;  
    ...  
    CPU.r0 = TCB[tNew].regs.r0;  
    CPU.sp = TCB[tNew].regs.sp;  
    CPU.retpc = TCB[tNew].regs.retpc;  
    return; /* Return to CPU.retpc */  
}
```

## Switch Details (continued)

---


- What if you make a mistake in implementing switch?
  - Suppose you forget to save/restore register 32
  - Get intermittent failures depending on when context switch occurred and whether new thread uses register 32
  - System will give wrong result without warning
- Can you devise an exhaustive test to test switch code?
  - No! Too many combinations and inter-leavings
- Cautionary tale:
  - For speed, Topaz kernel saved one instruction in switch()
  - Carefully documented! Only works as long as kernel size < 1MB
  - What happened?
    - » Time passed; people forgot
    - » Later, they added features to kernel (no one removes features!)
    - » Very weird behavior started happening
  - Moral of story: Design for simplicity




# How expensive is context switching?

- Switching between threads in same process similar to switching between threads in different processes, but *much cheaper*:
  - No need to change address space
- Some numbers from Linux:
  - Frequency of context switch: 10-100ms
  - Switching between processes: 3-4  $\mu$ sec.
  - Switching between threads: 100 ns
- Even cheaper: switch threads (using “yield”) in user-space!


*What we are talking about  
in Today's lecture*



Simple One-to-One




Many-to-One



Many-to-Many

# What happens when thread blocks on I/O?

---

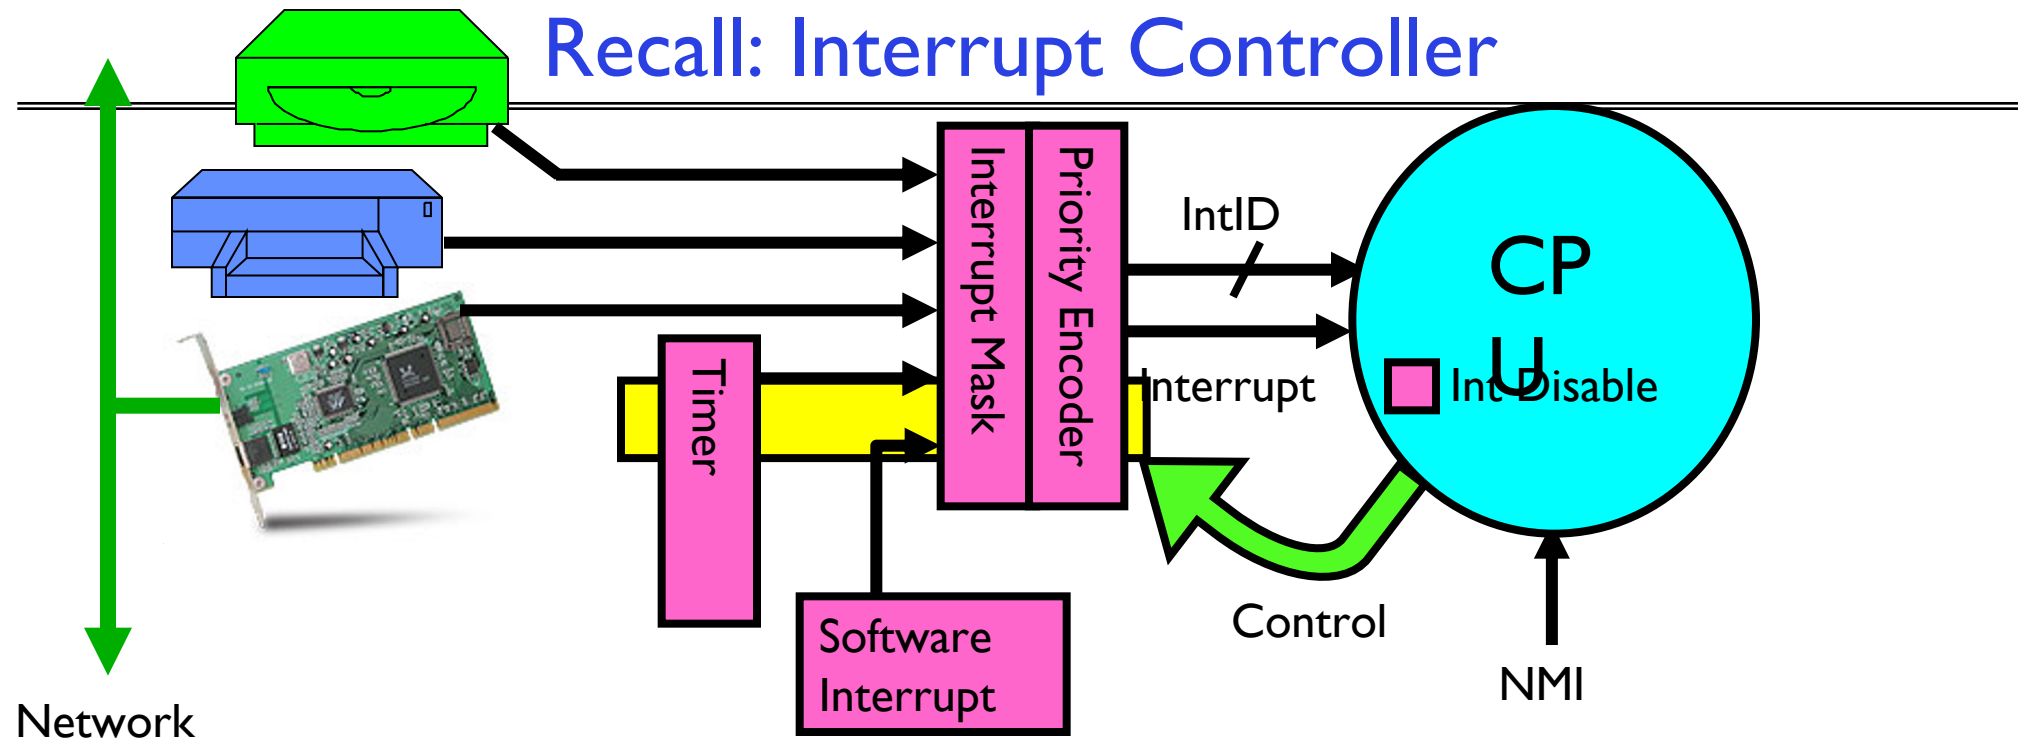


- What happens when a thread requests a block of data from the file system?
  - User code invokes a system call
  - Read operation is initiated
  - Run new thread/switch
- Thread communication similar
  - Wait for Signal/Join
  - Networking

# External Events

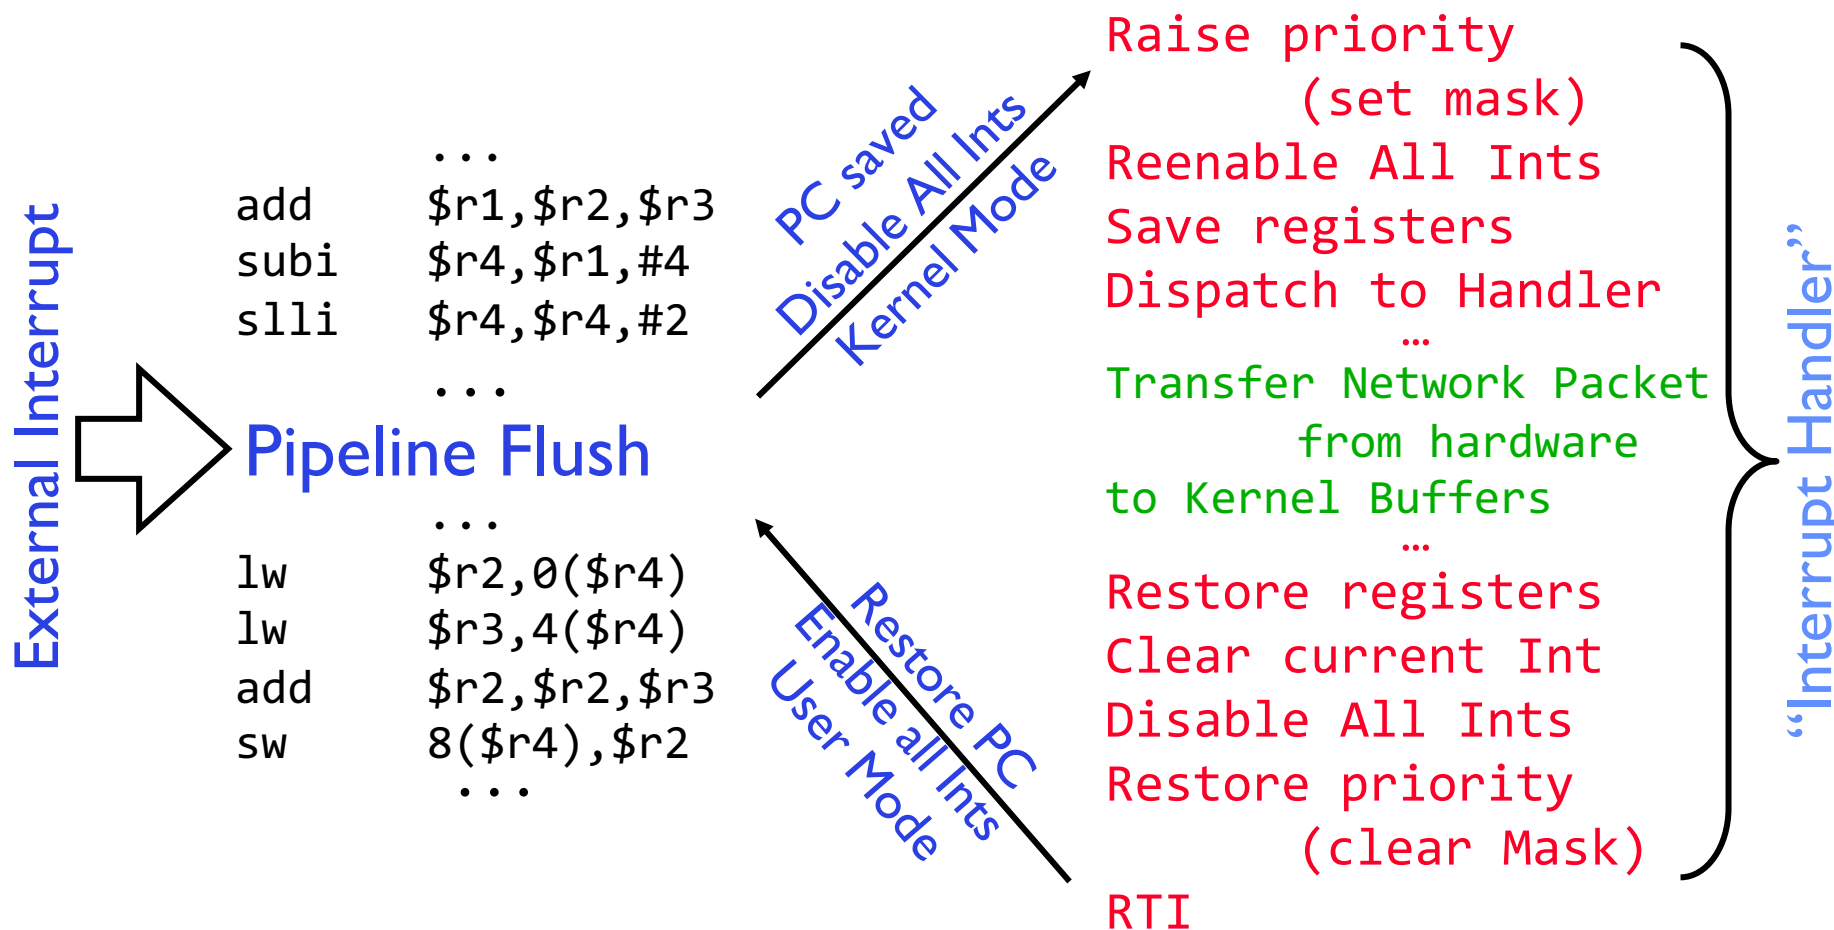
---

- What happens if thread never does any I/O, never waits, and never yields control?
  - Could the ComputePI program grab all resources and never release the processor?
    - » What if it didn't print to console?
  - Must find way that dispatcher can regain control!
- Answer: utilize external events
  - Interrupts: signals from hardware or software that stop the running code and jump to kernel
  - Timer: like an alarm clock that goes off every few milliseconds
- If we make sure that external events occur frequently enough, can ensure dispatcher runs



- Interrupts invoked with interrupt lines from devices
- Interrupt controller chooses interrupt request to honor
  - Interrupt identity specified with ID line
  - Mask enables/disables interrupts
  - Priority encoder picks highest enabled interrupt
  - Software Interrupt Set/Cleared by Software
- CPU can disable all interrupts with internal flag
- Non-Maskable Interrupt line (NMI) can't be disabled

# Example: Network Interrupt




- An interrupt is a hardware-invoked context switch
  - No separate step to choose what to run next
  - Always run the interrupt handler immediately

# Use of Timer Interrupt to Return Control

---

- Solution to our dispatcher problem
  - Use the timer interrupt to force scheduling decisions



- Timer Interrupt routine:

```
TimerInterrupt() {  
    DoPeriodicHouseKeeping();  
    run_new_thread();  
}
```

# ThreadFork(): Create a New Thread

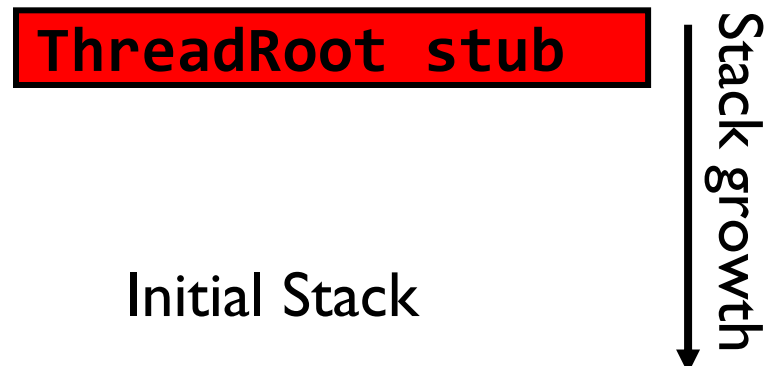
---

- ThreadFork() is a user-level procedure that creates a new thread and places it on ready queue
- Arguments to ThreadFork()
  - Pointer to application routine (fcnPtr)
  - Pointer to array of arguments (fcnArgPtr)
  - Size of stack to allocate
- Implementation
  - Sanity check arguments
  - Enter Kernel-mode and Sanity Check arguments again
  - Allocate new Stack and TCB
  - Initialize TCB and place on ready list (Runnable)

# How do we initialize TCB and Stack?

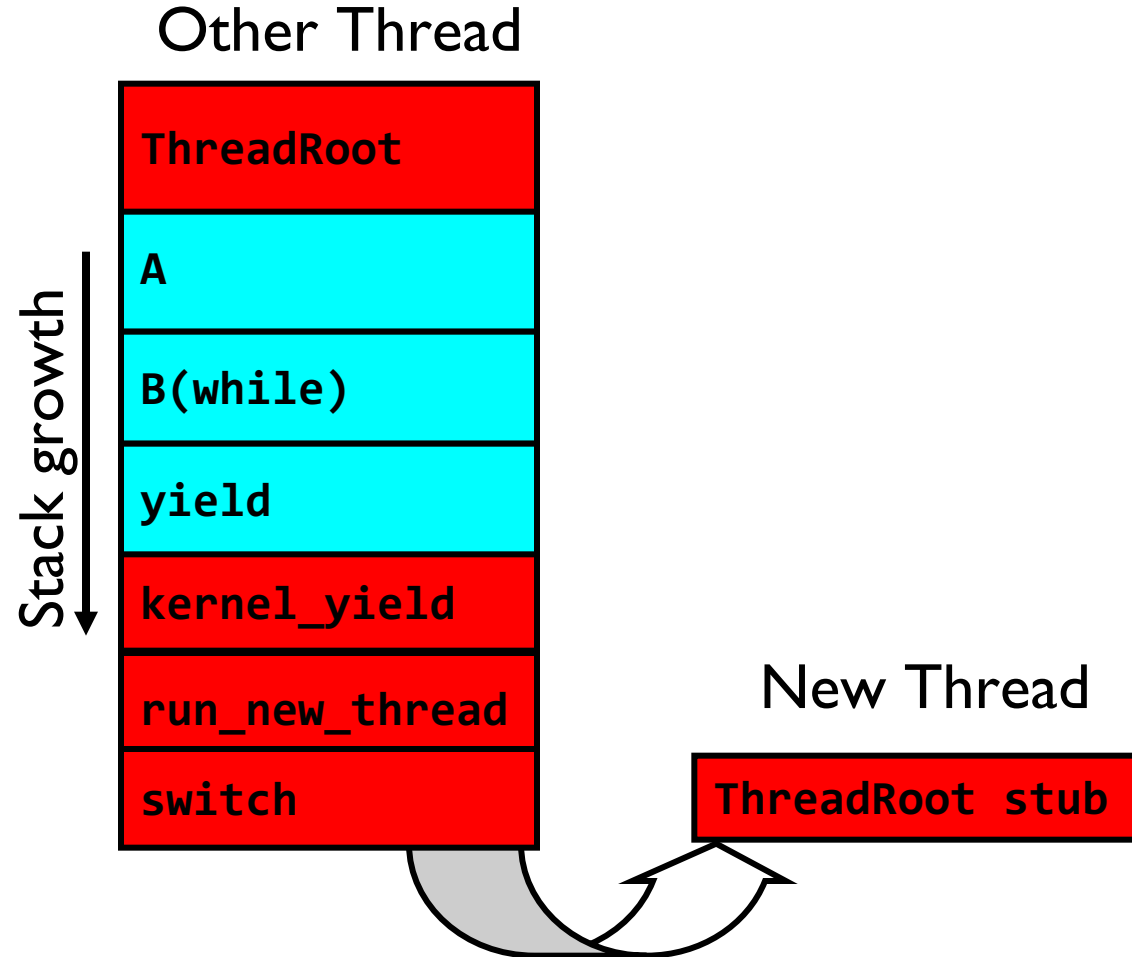
---

- Initialize Register fields of TCB
  - Stack pointer made to point at stack
  - PC return address  $\Rightarrow$  OS (asm) routine ThreadRoot()
  - Two arg registers (a0 and a1) initialized to fcnPtr and fcnArgPtr, respectively
- Initialize stack data?
  - Minimal initialization  $\Rightarrow$  setup return to go to beginning of ThreadRoot()
    - » Important part of stack frame is in registers for RISC-V (ra)
    - » X86: need to push a return address on stack
  - Think of stack frame as just before body of ThreadRoot() really gets started



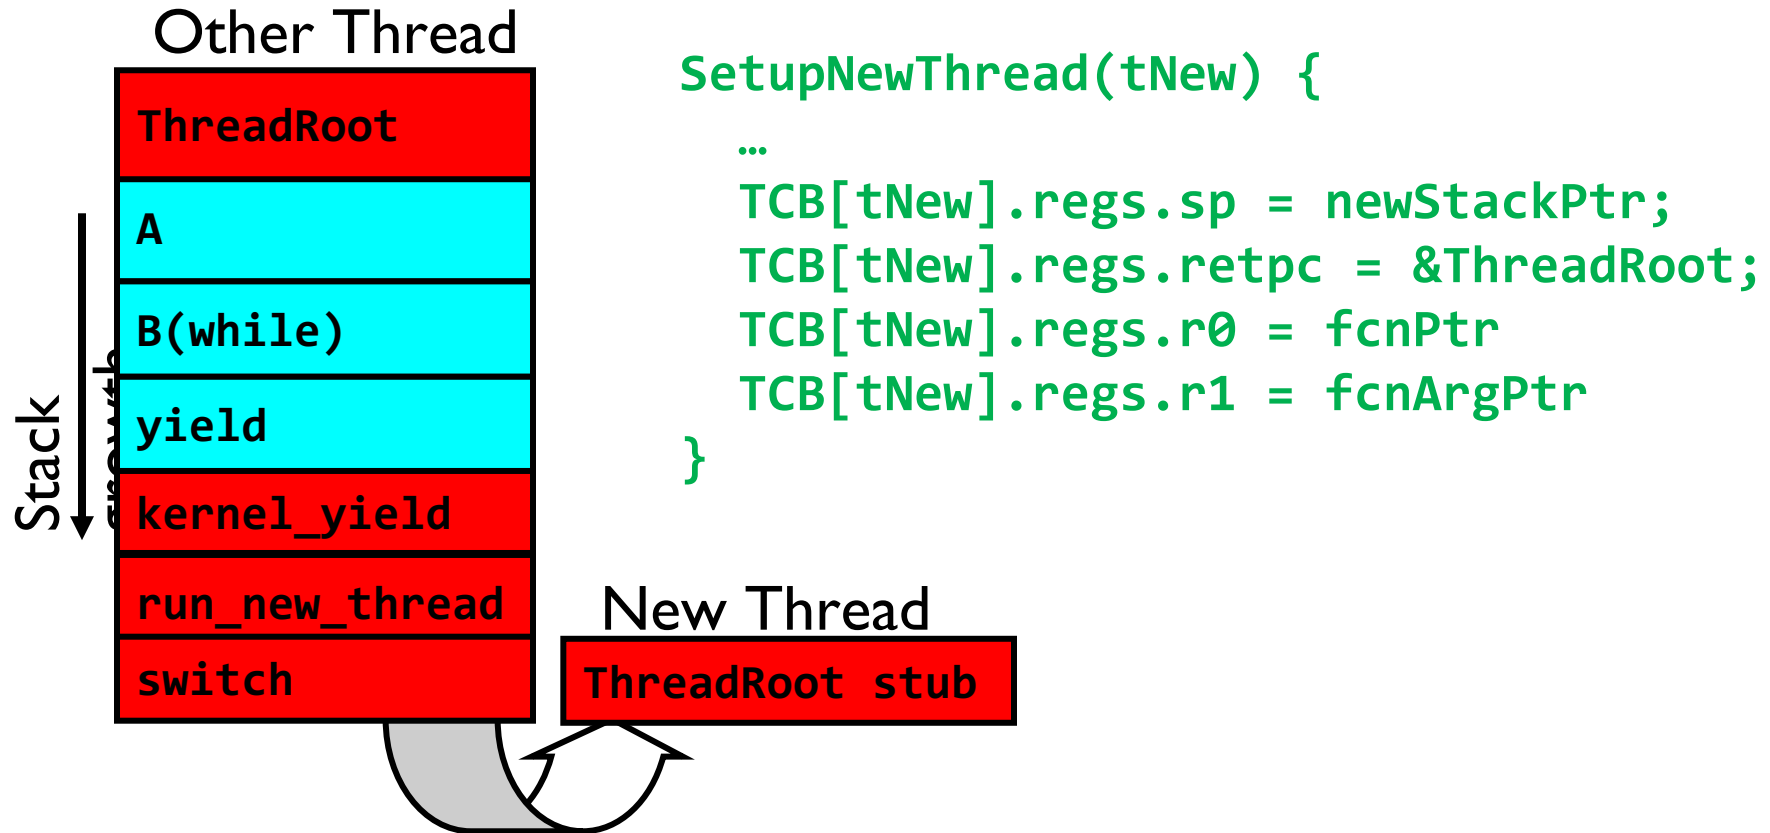


# How does Thread get started?



- Eventually, `run_new_thread()` will select this TCB and return into beginning of `ThreadRoot()`
  - This really starts the new thread

# How does a thread get started?




- How do we make a **new** thread?
  - Setup TCB/kernel thread to point at new user stack and ThreadRoot code
  - Put pointers to start function and args in registers or top of stack
    - » This depends heavily on the calling convention (i.e. RISC-V vs x86)
- Eventually, run\_new\_thread() will select this TCB and return into beginning of ThreadRoot()
  - This really starts the new thread

# What does ThreadRoot() look like?

- ThreadRoot() is the root for the thread routine:

```
ThreadRoot(fcnPTR, fcnArgPtr) {  
    DoStartupHousekeeping();  
    UserModeSwitch(); /* enter user mode */  
    Call fcnPtr(fcnArgPtr);  
    ThreadFinish();  
}
```

- Startup Housekeeping
  - Includes things like recording start time of thread
  - Other statistics
- Stack will grow and shrink with execution of thread
- Final return from thread returns into ThreadRoot() which calls ThreadFinish()
  - ThreadFinish() wake up sleeping threads



# Threads vs Address Spaces: Options

# threads Per AS:	# of addr spaces:	One	Many
One	MS/DOS, early Macintosh	Traditional UNIX	
Many	Embedded systems (Geoworks, VxWorks, JavaOS,etc) JavaOS, Pilot(PC)	Mach, OS/2, Linux Windows 10 Win NT to XP, Solaris, HP-UX, OS X	

- Most operating systems have either
  - One or many address spaces
  - One or many threads per address space

# Administrivia

---

- Midterm Thursday 10/3
  - Closed book, but one page of *handwritten* notes, both sides
  - No class on day of midterm
  - 7-9PM
- Project I Design Document due next Thursday 9/26
  - No extensions of any sort on design documents!!!
- Project I Design reviews upcoming
  - High-level discussion of your approach
    - » What will you modify?
    - » What algorithm will you use?
    - » How will things be linked together, etc.
    - » Do not need final design (complete with all semicolons!)
  - You will be asked about testing
    - » Understand testing framework
    - » Are there things you are doing that are not tested by tests we give you?
- Do your own work!
  - Please do not try to find solutions from previous terms
  - We will be on the look out for anyone doing this...today

# Goals for Rest of Today


---

- Challenges and Pitfalls of Concurrency
- Synchronization Operations/Critical Sections
- How to build a lock?
- Atomic Operations




# Concurrency vs Parallelism

- Multithreading: Multiple threads per Process (A *programming strategy*)
- Multiplexing: Sharing a single resource (such as a core) among multiple threads
- What does it mean to run two threads “concurrently” (*regardless of process*)?
  - Scheduler is free to run threads in any order and interleaving: FIFO, Random, ...
  - Unless synchronization is involved, multiple threads are concurrent!
  - Assume: if scheduler can produce the worst possible interleaving, IT WILL!




- What does it mean to run two threads “in parallel” (*regardless of process*)?
  - Threads are *actually running* at the same time
  - Parallel  $\Rightarrow$  Concurrent but Concurrent  $\nRightarrow$  Parallel



# ATM Bank Server

---



- ATM server problem:
  - Service a set of requests
  - Do so without corrupting database
  - Don't hand out too much money



# ATM bank server example

---

- Suppose we wanted to implement a server process to handle requests from an ATM network:

```
BankServer() {
    while (TRUE) {
        ReceiveRequest(&op, &acctId, &amount);
        ProcessRequest(op, acctId, amount);
    }
}

ProcessRequest(op, acctId, amount) {
    if (op == deposit) Deposit(acctId, amount);
    else if ...
}

Deposit(acctId, amount) {
    acct = GetAccount(acctId); /* may use disk I/O */
    acct->balance += amount;
    StoreAccount(acct); /* Involves disk I/O */
}
```

- How could we speed this up?
  - More than one request being processed at once
  - Event driven (overlap computation and I/O)
  - Multiple threads (multi-proc, or overlap comp and I/O)

# Event Driven Version of ATM server

---

- Suppose we only had one CPU
  - Still like to overlap I/O with computation
  - Without threads, we would have to rewrite in event-driven style

- Example

```
BankServer() {  
    while(TRUE) {  
        event = WaitForNextEvent();  
        if (event == ATMRequest)  
            StartOnRequest();  
        else if (event == AcctAvail)  
            ContinueRequest();  
        else if (event == AcctStored)  
            FinishRequest();  
    }  
}
```

- This technique is used for graphical programming

- Complication:
  - What if we missed a blocking I/O step?
  - What if we have to split code into hundreds of pieces which could be blocking?

# Can Threads (in same Process) Make This Easier?

---

- Threads yield overlapped I/O and computation without “deconstructing” code into non-blocking fragments
  - One thread per request
- Requests proceeds to completion, blocking as required:

```
Deposit(acctId, amount) {  
    acct = GetAccount(actId); /* May use disk I/O */  
    acct->balance += amount;  
    StoreAccount(acct);      /* Involves disk I/O */  
}
```

- Unfortunately, shared state can get corrupted:

<u>Thread 1</u>	<u>Thread 2</u>
load r1, acct->balance	
	load r1, acct->balance
	add r1, amount2
	store r1, acct->balance
add r1, amount1	
store r1, acct->balance	

# Problem is at the Lowest Level

---

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A      Thread B

$x = 1; y = 2;$

- However, what about (Initially,  $y = 12$ ):

Thread A      Thread B

$x = 1; y = 2;$

$x = y + 1; \quad y = y * 2;$

- What are the possible values of  $x$ ?
- Or, what are the possible values of  $x$  below?

Thread A      Thread B

$x = 1; x = 2;$

- $x$  could be 1 or 2 (non-deterministic!)
- Could even be 3 for serial processors:
  - » Thread A writes 0001, B writes 0010 → scheduling order ABABABBA yields 3!

# Atomic Operations

---

- To understand a concurrent program, we need to know what the underlying indivisible operations are!
- **Atomic Operation**: an operation that always runs to completion or not at all
  - It is *indivisible*: it cannot be stopped in the middle and state cannot be modified by someone else in the middle
  - Fundamental building block – if no atomic operations, then have no way for threads to work together
- On most machines, memory references and assignments (i.e. loads and stores) of words are atomic
  - Consequently – weird example that produces “3” on previous slide can’t happen
- Many instructions are not atomic
  - Double-precision floating point store often not atomic
  - VAX and IBM 360 had an instruction to copy a whole array

# Another Concurrent Program Example

---

- Two threads, A and B, compete with each other
  - One tries to increment a shared counter
  - The other tries to decrement the counter

<u>Thread A</u>	<u>Thread B</u>
<code>i = 0;</code>	<code>i = 0;</code>
<code>while (i &lt; 10)</code>	<code>while (i &gt; -10)</code>
<code>    i = i + 1;</code>	<code>    i = i - 1;</code>
<code>printf("A wins!");</code>	<code>printf("B wins!");</code>

- Assume that memory loads and stores are atomic, but incrementing and decrementing are *not* atomic
- Who wins? Could be either
- Is it guaranteed that someone wins? Why or why not?
- What if both threads have their own CPU running at same speed? Is it guaranteed that it goes on forever?

# Hand Simulation Multiprocessor Example

---

- Inner loop looks like this:

	<u>Thread A</u>	<u>Thread B</u>
r1=0	load r1, M[i]	
		r1=0 load r1, M[i]
r1=1	add r1, r1, 1	
		r1=-1 sub r1, r1, 1
M[i]=1	store r1, M[i]	
	M[i]=-1	store r1, M[i]

- **Hand Simulation:**
  - And we're off. A gets off to an early start
  - B says “hmp, better go fast” and tries really hard
  - A goes ahead and writes “1”
  - B goes and writes “-1”
  - A says “HUH??? I could have sworn I put a 1 there”
- Could this happen on a uniprocessor? With Hyperthreads?
  - Yes! Unlikely, but if you are depending on it not happening, it will and your system will break...

# Definitions

---

- **Synchronization**: using atomic operations to ensure cooperation between threads
  - For now, only loads and stores are atomic
  - We are going to show that its hard to build anything useful with only reads and writes
- **Mutual Exclusion**: ensuring that only one thread does a particular thing at a time
  - One thread *excludes* the other while doing its task
- **Critical Section**: piece of code that only one thread can execute at once. Only one thread at a time will get into this section of code
  - Critical section is the result of mutual exclusion
  - Critical section and mutual exclusion are two ways of describing the same thing



# Locks

---



- **Lock**: prevents someone from doing something
  - **Lock()** before entering critical section and before accessing shared data
  - **Unlock()** when leaving, after accessing shared data
  - **Wait** if locked
    - » Important idea: all synchronization involves waiting
- Locks need to be allocated and initialized:
  - `structure Lock mylock`      or      `pthread_mutex_t mylock;`
  - `lock_init(&mylock)`      or      `mylock = PTHREAD_MUTEX_INITIALIZER;`
- Locks provide two **atomic** operations:
  - **acquire(&mylock)** – wait until lock is free; then mark it as busy
    - » After this returns, we say the calling thread *holds* the lock
  - **release(&mylock)** – mark lock as free
    - » Should only be called by a thread that currently holds the lock
    - » After this returns, the calling thread no longer holds the lock

## Fix banking problem with Locks!

- Identify critical sections (atomic instruction sequences) and add locking:

**Code Snippet:**

```

Deposit(acctId, amount) {
    acquire(&mylock)           // Wait if someone else in critical
    section!
    acct = GetAccount(actId);
    acct->balance += amount;
    StoreAccount(acct);
    release(&mylock)           // Release someone into critical
}

```

**Thread Timeline:**

- Thread A:** Starts, enters the critical section (red box), and releases the lock.
- Thread B:** Starts, enters the critical section (red box), and releases the lock.
- Thread C:** Starts, enters the critical section (red box), and releases the lock.

The diagram shows that only one thread is in the critical section at any given time, ensuring mutual exclusion.

- Must use SAME lock (mylock) with all of the methods (Withdraw, etc...)

# Correctness Requirements

- Threaded programs must work for all interleavings of thread instruction sequences
  - Cooperating threads inherently non-deterministic and non-reproducible
  - Really hard to debug unless carefully designed!
- Example: Therac-25
  - Machine for radiation therapy
    - » Software control of electron accelerator and electron beam/Xray production
    - » Software control of dosage
  - Software errors caused the death of several patients
    - » A series of race conditions on shared variables and poor software design
    - » “They determined that data entry speed during editing was the key factor in producing the error condition: If the prescription data was edited at a fast pace, the overdose occurred.”




Figure 1. Typical Therac-25 facility

# Conclusion

---

- Every thread has both a user and kernel stack
  - Showed more details about context-switching mechanisms
- Concurrent threads introduce problems when accessing shared data
  - Programs must be insensitive to arbitrary interleavings
  - Without careful design, shared variables can become completely inconsistent
- Important concept: **Atomic Operations**
  - An operation that runs to completion or not at all
  - These are the primitives on which to construct various synchronization primitives
- Introduced the Lock API: **acquire()** and **release()**
  - Next time: How do we make a lock?